

Gunrock: A High Performance Graph Processing Library on the GPU

June 4, 2015

Yangzihao Wang
University of California, Davis

PROBLEM

High-performance large-scale graph analytic

PROBLEM

**High-performance large-scale graph analytic
on GPUs**

Why Using GPUs for Graph Analytic?

- **Graphs are ubiquitous**
- **Data size is becoming very large**
- **Graph analytic systems demand more performance**

Large-scale Graph Analytic Is Difficult

- Irregularity of data access and control flow limits performance and scalability
- GPU programming is complex

Related Work

- **Single-node CPU-based systems: Boost Graph Library**
- **Distributed CPU-based systems: PowerGraph, Pregel**
- **Specialized GPU algorithms and GPU-based systems**

Related Work

- Single-node CPU-based systems: Boost Graph Library
- Distributed CPU-based systems: PowerGraph, Pregel
- Specialized GPU algorithms and GPU-based systems

Can we do better? (in terms of both performance and expressiveness)

IDEA: Performance AND expressiveness

- **Performance: Integrating high performance GPU computing primitives and optimizations into the core.**
- **Expressiveness: A data-centric abstraction designed specifically for the GPU**

IDEA: Bulk-Synchronous Programming and Data-Centric

- **Graph algorithms as iterative convergent processes**
 - A series of bulk synchronous steps
 - Large amount of parallelism within each step
- **Manipulate frontiers**
 - Generating/reorganizing frontier in parallel
 - Computing on frontier in parallel

Expressiveness: Gunrock's Data-Centric Abstraction


Gunrock's Key Abstraction Is **FRONTIER**

Most graph algorithms have two major operations:

- **Traverse:** moving in the graph and generating new frontier
- **Compute:** doing computation on frontier


Gunrock's Traversal Step

- **Advance:** visiting the neighbors of the current frontier
- **Filter:** choosing a subset of the current frontier




Gunrock's Traversal Step

- **Advance:** visiting the neighbors of the current frontier
- **Filter:** choosing a subset of the current frontier




Gunrock's Traversal Step

- **Advance:** visiting the neighbors of the current frontier
- **Filter:** choosing a subset of the current frontier




Gunrock's Traversal Step

- **Advance:** visiting the neighbors of the current frontier
- **Filter:** choosing a subset of the current frontier



Gunrock's Traversal Step


- **Advance:** visiting the neighbors of the current frontier
- **Filter:** choosing a subset of the current frontier



Gunrock's Compute Step

Functors that apply to {edges, vertices}

- “cond” functor: returns a boolean value
- “apply” functor: performs a computation




Graph Primitives in Gunrock (In only Three Files)

- **Problem:** Initialize the graph data and frontier
- **Enactor:** GPU kernel entry function which defines a series of operations on frontier
- **Functor:** User-specified per-node/per-edge computation on frontier

Performance: Generalized Optimization Strategies

Graph Data Representation: CSR

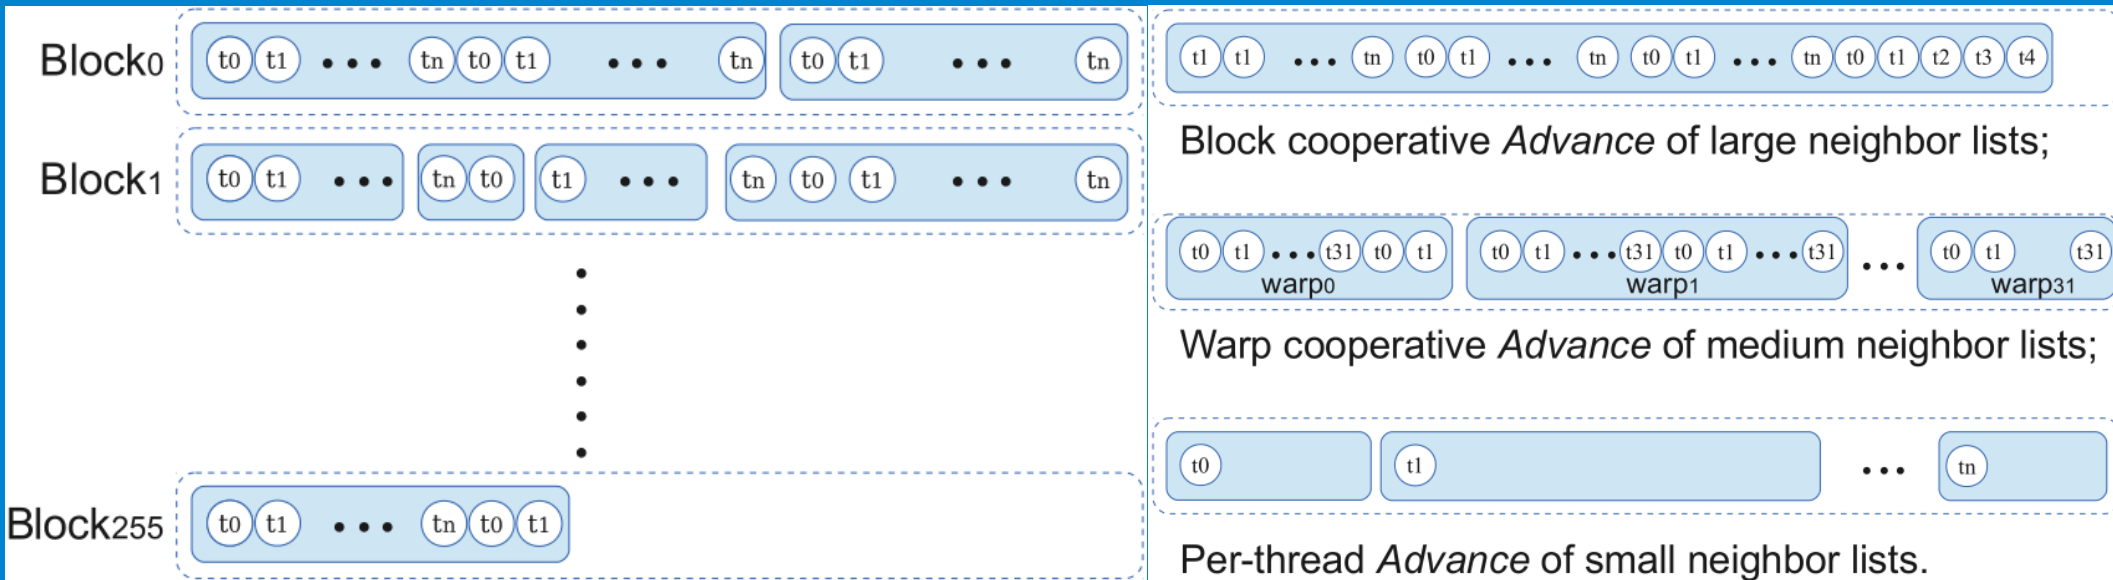


Workload Mapping and Load-balancing

- **Naive method: Let one thread handle the neighbor list of one vertex**
- **Problem: Highly uneven distribution of node degrees in scale-free graphs**

Need load balancing strategy!

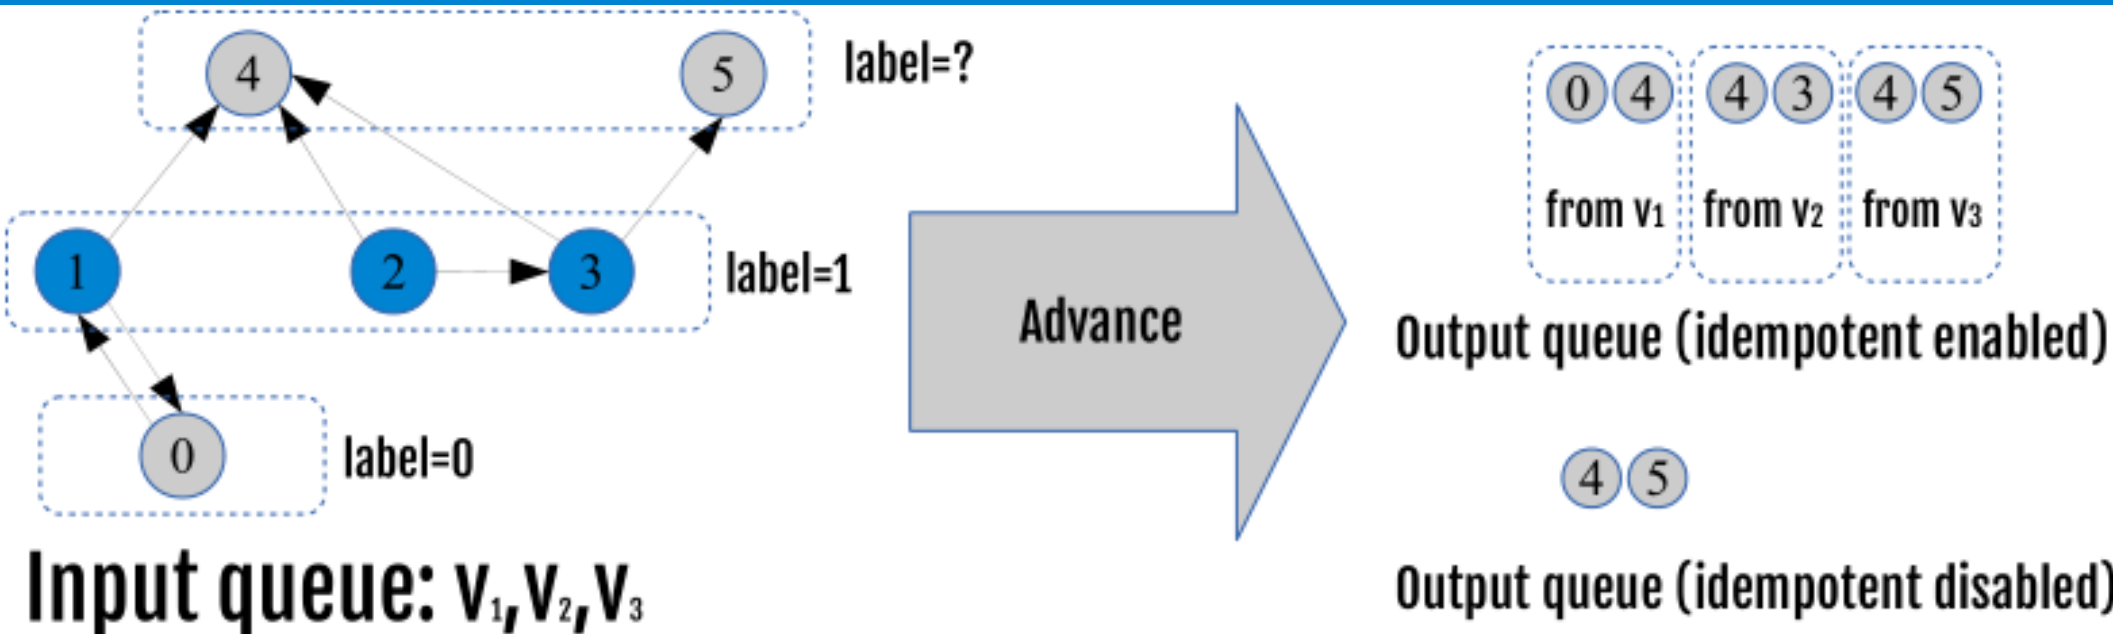
Workload Mapping and Load-balancing



- Tradeoff between extra processing and load balancing
- A worthwhile extra effort: 2x–20x speedup over non-load balancing library (Medusa v1)

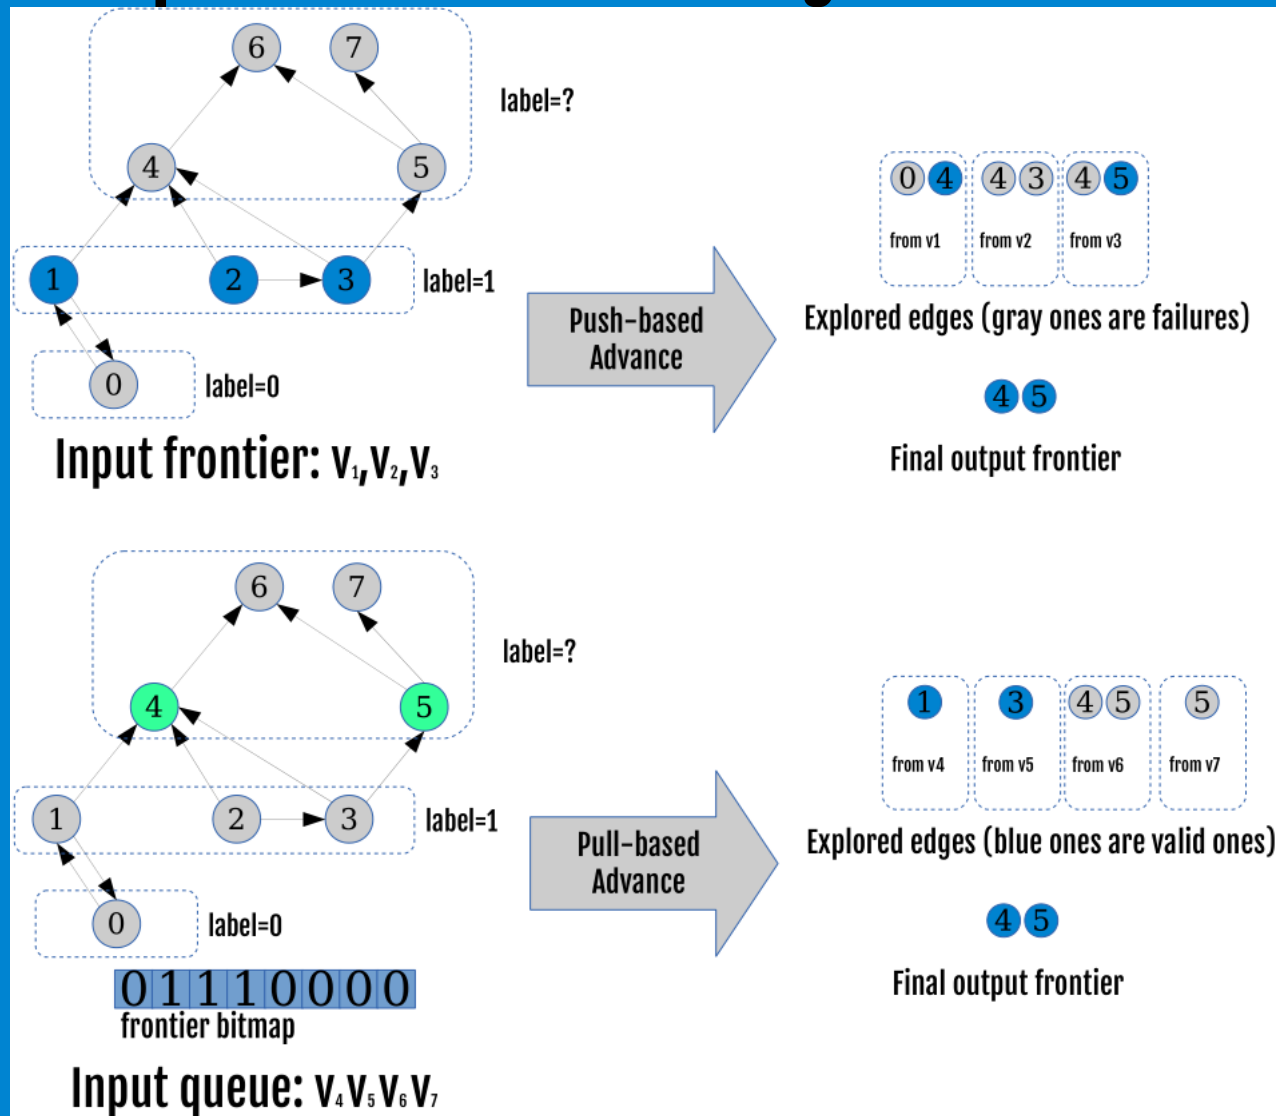
Data-Centric Abstraction Enables Optimizations

Idempotent operations (frontier reorganization)




Data-Centric Abstraction Enables Optimizations

Pull vs. push operations (frontier generation)




Data-Centric Abstraction Enables Optimizations

Priority Queue (frontier reorganization)



Results, Conclusion, and Future Work

Performance Against Other Graph Processing Systems



Expressiveness and Usability

Currently have over 10 graph primitives

- **Traversal-based, Node-ranking, Global (connected component, MST)**
- **LOC under 300 for each primitive**

Working on more graph primitives

- **Graph coloring, Maximal Independent Set**
- **Community Detection**
- **Subgraph Matching**

Future Works

Operations

- Gather reduce
- Global indicator
- Operations on two frontier sets (intersection, join)
- Sampling frontier with different distributions
- Form supervertex

Future Works

Primitives

- Stochastic Gradient Descent
- Markov-Chain Monte Carlo

Compiler

- Integration with TinkerPop (Popgun!)
- Explorer kernel fusion

Conclusions

High-level Abstraction is essential for GPUs to make an impact in graph analytic.

Gunrock's data-centric, frontier-focused abstraction has good balance between expressiveness and performance

Questions?

Gunrock Team == Yangzihao Wang, Yuechao Pan, Yuduo Wu, Carl Yang, Leyuan Wang, and our advisor: John D. Owens!